

Fundamentals of Stream Processing with Apache Beam (incubating)



Frances Perry & Tyler Akidau

@francesjperry, @takidau

Apache Beam Committers & Google Engineers

Agenda

1

Infinite, Out-of-Order Data Sets

2

What, Where, When, How

3

Reasons This is Awesome


4

Apache Beam (incubating)




Infinite, Out-of-Order Data Sets

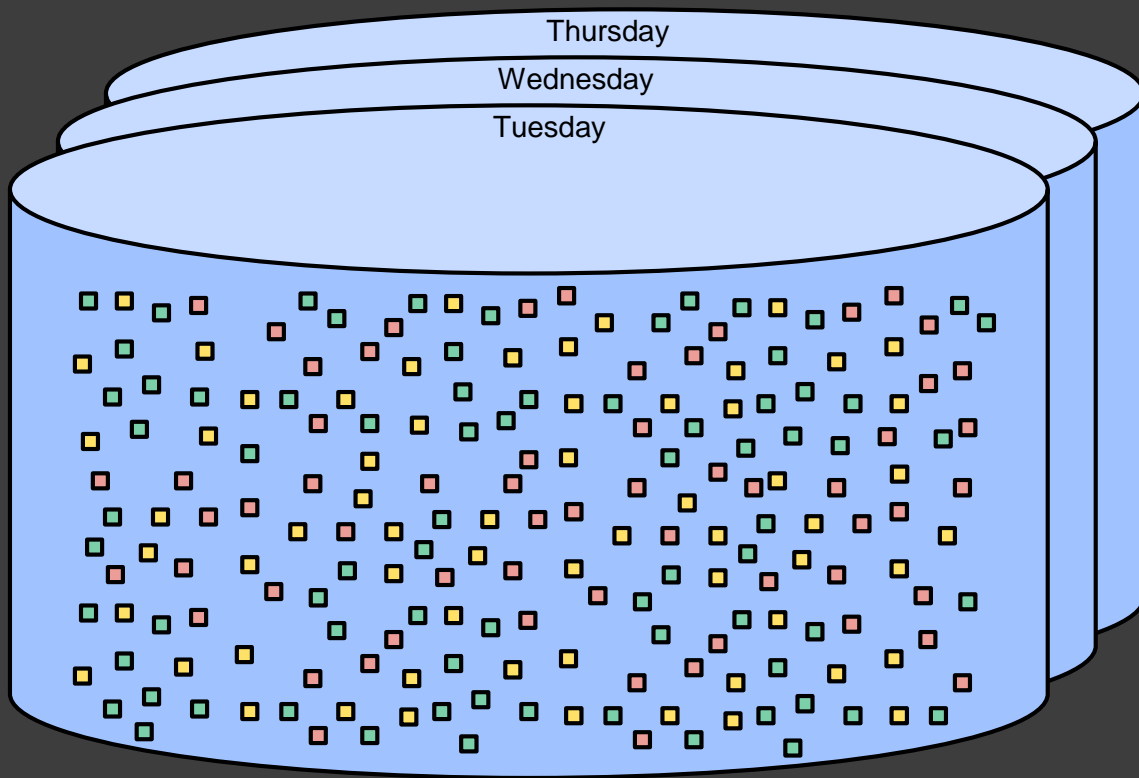
Data...




...can be big...



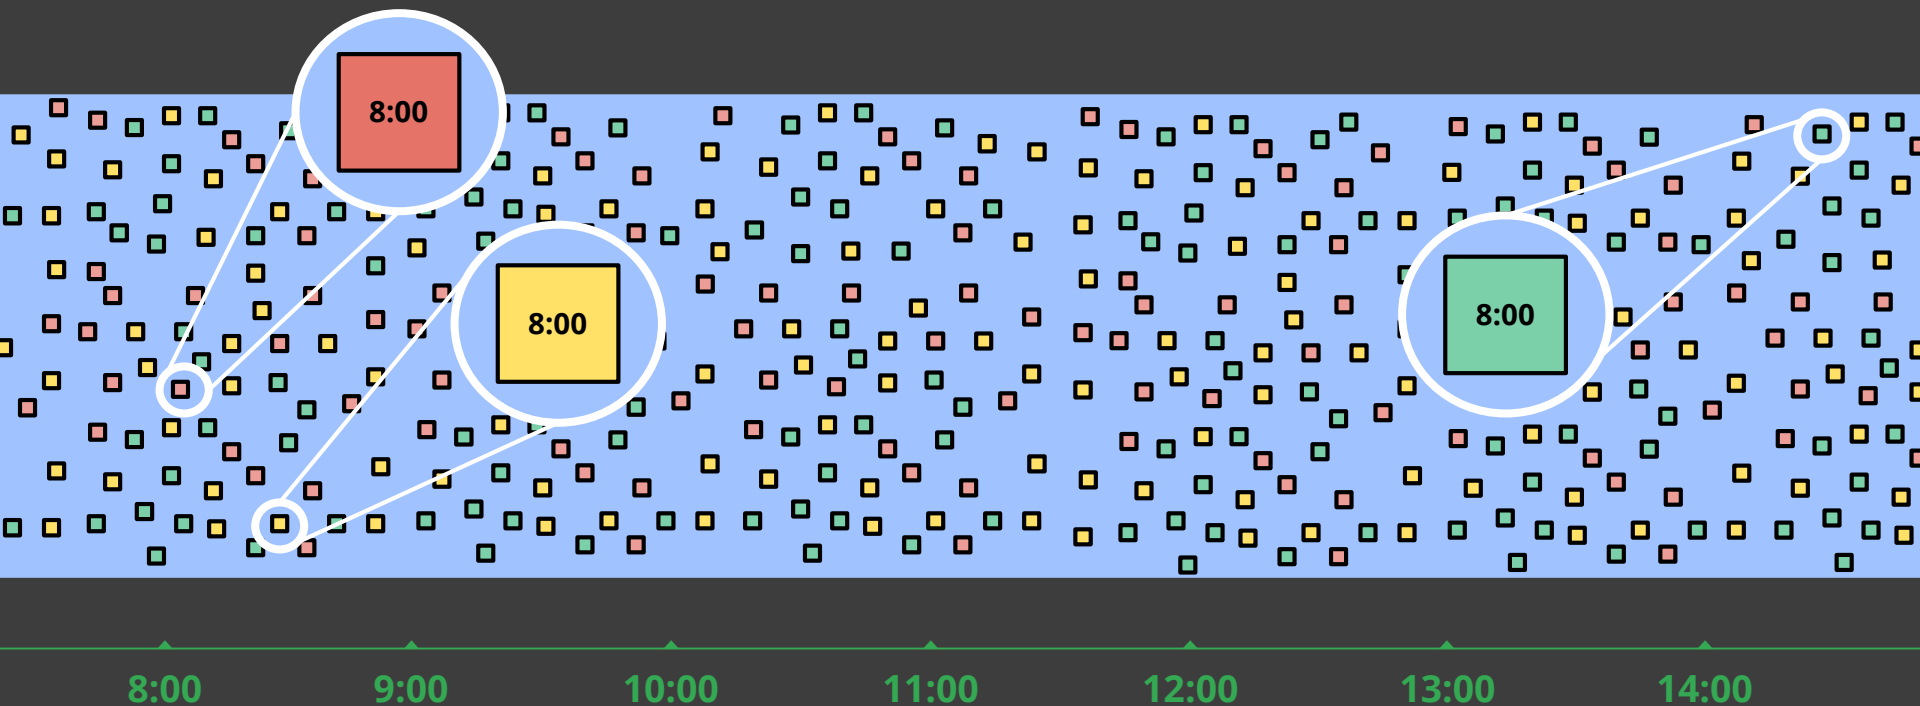
...really, really big...




... maybe infinitely big...




... with unknown delays.




Element-wise transformations




Aggregating via Processing-Time Windows




Aggregating via Event-Time Windows



Formalizing Event-Time Skew



Formalizing Event-Time Skew



Watermarks describe event time progress.

"No timestamp earlier than the watermark will be seen"

Often heuristic-based.

Too Slow? Results are *delayed*.
Too Fast? Some data is *late*.



What, Where, When, How


What are you computing?

Where in event time?


When in processing time?

How do refinements relate?


What are you computing?



Element-Wise



Aggregating



Composite


What: Computing Integer Sums

```
// Collection of raw log lines
PCollection<String> raw = IO.read(...);


// Element-wise transformation into team/score pairs
PCollection<KV<String, Integer>> input =
    raw.apply(ParDo.of(new ParseFn()));

// Composite transformation containing an aggregation
PCollection<KV<String, Integer>> scores =
    input.apply(Sum.integersPerKey());
```

What: Computing Integer Sums




What: Computing Integer Sums



Where in event time?

Windowing divides data into event-time-based finite chunks.




Often required when doing aggregations over unbounded data.


Where: Fixed 2-minute Windows

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2))))
    .apply(Sum.integersPerKey());
```

Where: Fixed 2-minute Windows



When in processing time?

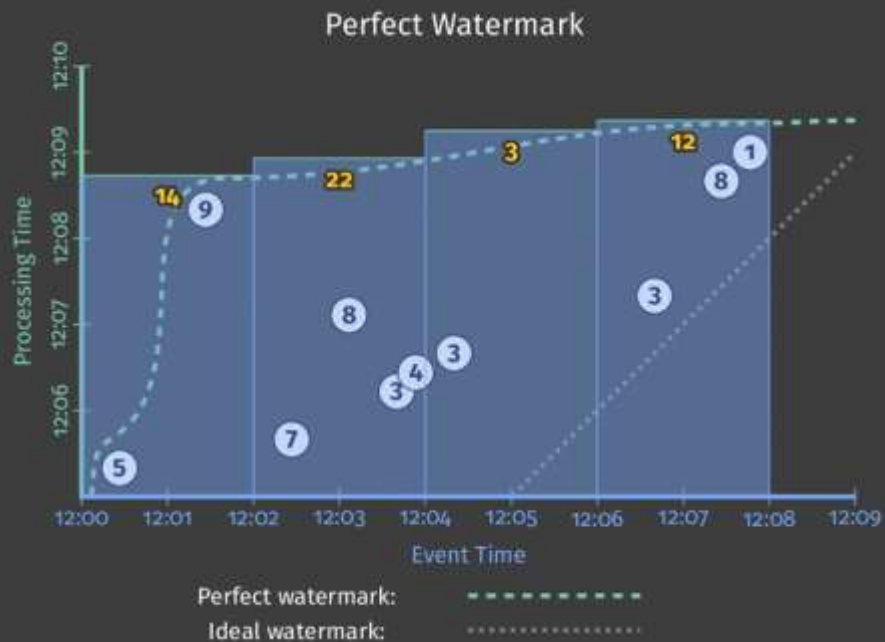


- Triggers control when results are emitted.
- Triggers are often relative to the watermark.

When: Triggering at the Watermark

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2))
        .triggering(AtWatermark())))
    .apply(Sum.integersPerKey());
```

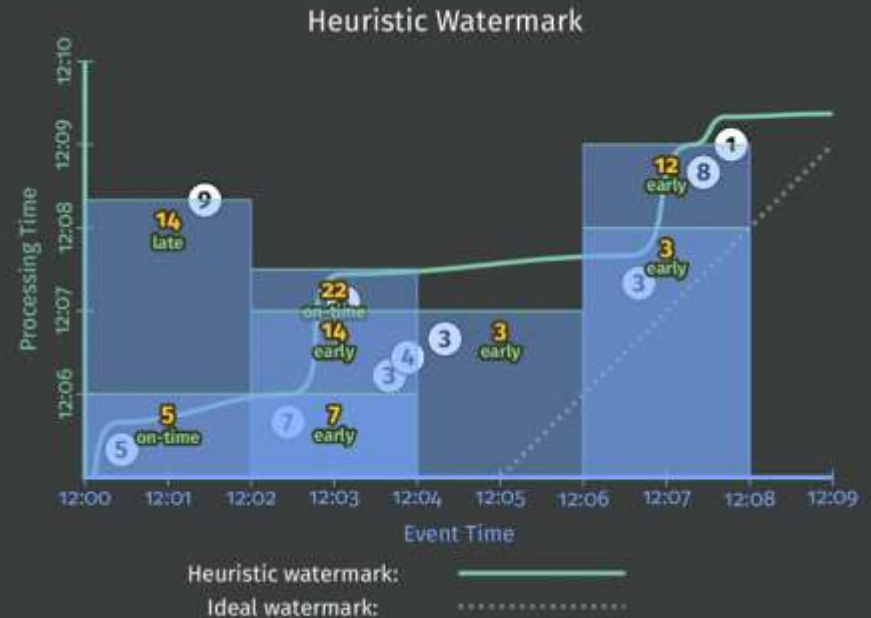
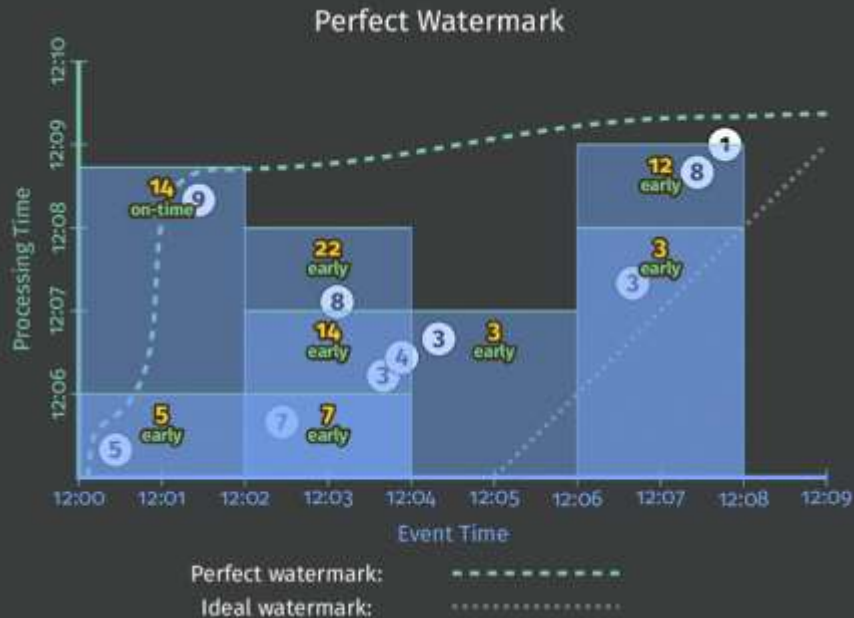

When: Triggering at the Watermark



When: Early and Late Firings

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2))
        .triggering(AtWatermark()
            .withEarlyFirings(AtPeriod(Minutes(1)))
            .withLateFirings(AtCount(1))))))
    .apply(Sum.integersPerKey());
```

When: Early and Late Firings



How do refinements relate?

- How should multiple outputs per window accumulate?
- Appropriate choice depends on consumer.


Firing	Elements	Discarding	Accumulating	Acc. & Retracting
Speculative	[3]	3	3	3
Watermark	[5, 1]	6	9	9, -3
Late	[2]	2	11	11, -9
<i>Last Observed</i>		2	11	11
<i>Total Observed</i>		11	23	11

(Accumulating & Retracting not yet implemented.)

How: Add Newest, Remove Previous

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2))
        .triggering(AtWatermark()
            .withEarlyFirings(AtPeriod(Minutes(1)))
            .withLateFirings(AtCount(1)))
        .accumulatingAndRetractingFiredPanels()))
    .apply(Sum.integersPerKey());
```

How: Add Newest, Remove Previous





Reasons This is Awesome

What / Where / When / How

Correctness

Power

Composability

Flexibility

Modularity



What / **Where** / **When** / **How**

Correctness

Power


Composability

Flexibility

Modularity



Distributed Systems are Distributed



Ideal watermark:

.....

Processing Time Results Differ



Event Time Results are Stable



What / Where / When / How

Correctness

Power

Composability

Flexibility

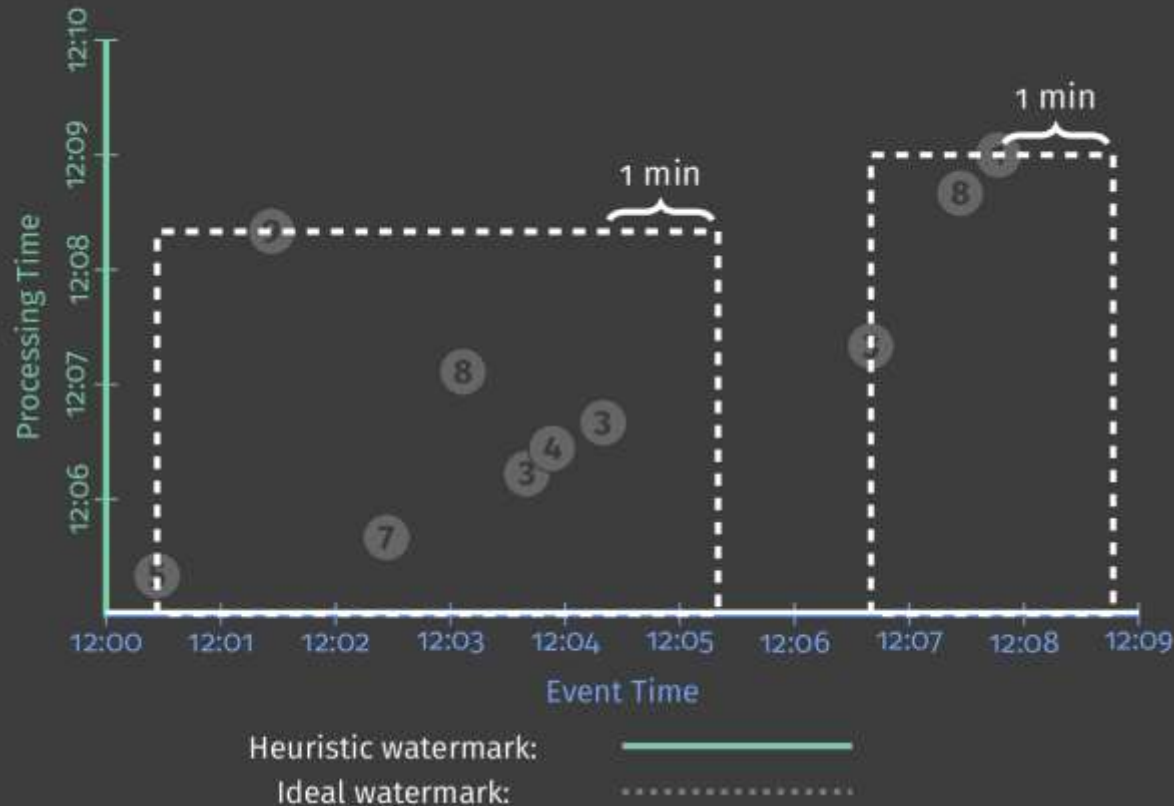
Modularity



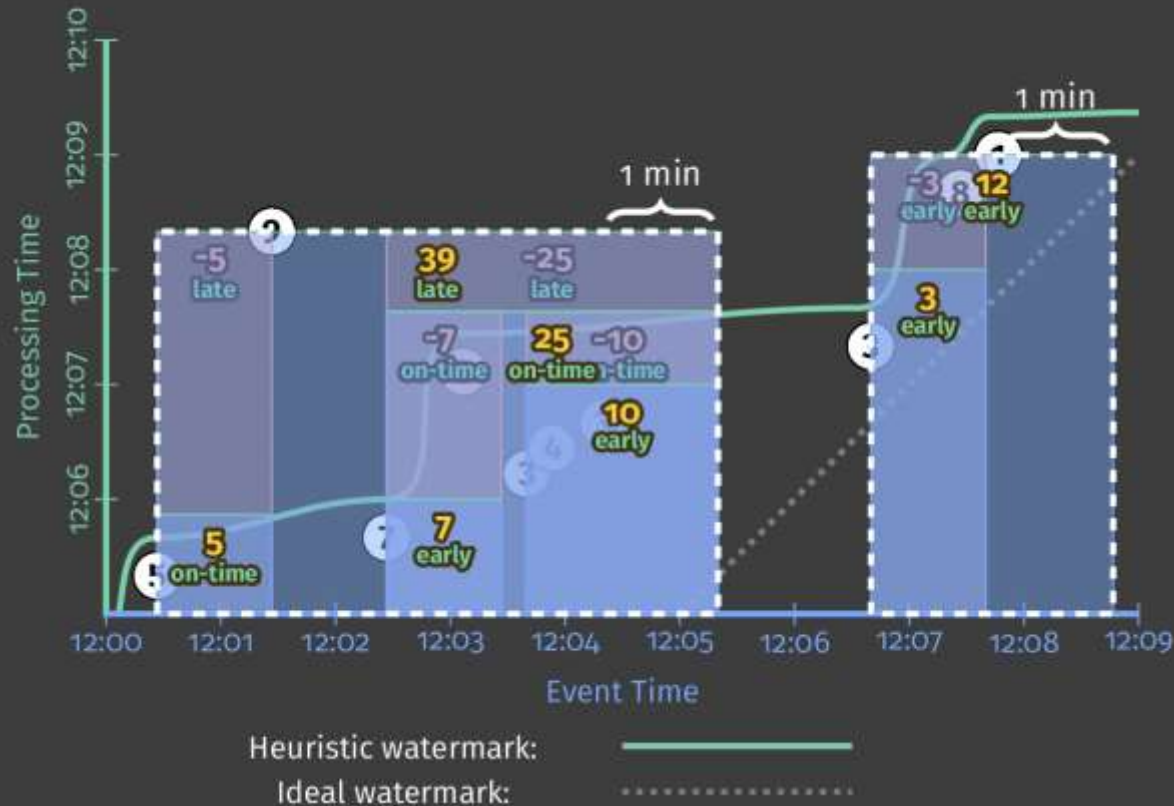
Sessions

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(Sessions.withGapDuration(Minutes(1)))
        .triggering(AtWatermark()
            .withEarlyFirings(AtPeriod(Minutes(1)))
            .withLateFirings(AtCount(1)))
        .accumulatingAndRetractingFiredPanels())
    .apply(Sum.integersPerKey());
```

Identifying Bursts of User Activity



Identifying Bursts of User Activity



What / Where / When / How

Correctness

Power


Composability

Flexibility


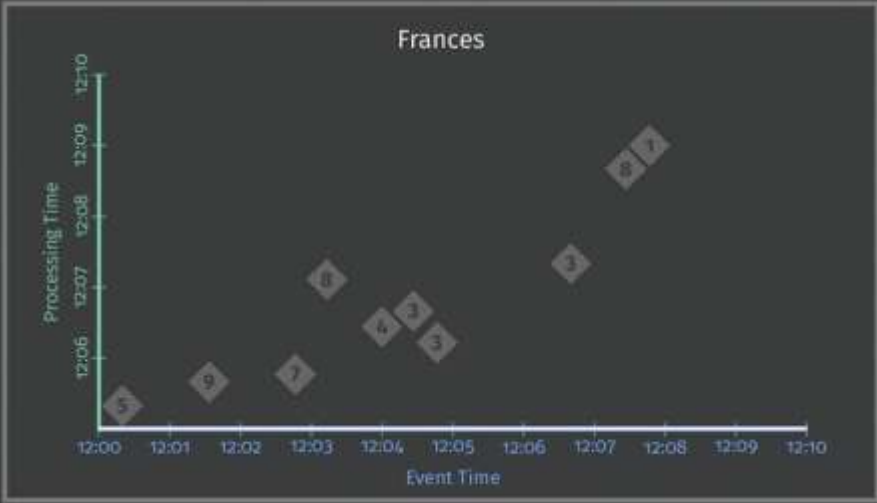
Modularity




Calculating Session Lengths




```
input
  .apply(Window.into(Sessions.withGapDuration(Minutes(1)))
    .trigger(AtWatermark())
    .discardingFiredPanes())
  .apply(CalculateWindowLength()));
```



Calculating the Average Session Length



```
input
    .apply(Window.into(Sessions.withGapDuration(Minutes(1))))
        .trigger(AtWatermark())
        .discardingFiredPanes())
    .apply(CalculateWindowLength());
```



```
.apply(Window.into(FixedWindows.of(Minutes(2)))
    .trigger(AtWatermark())
    .withEarlyFirings(AtPeriod(Minutes(1))))
    .accumulatingFiredPanes())
    .apply(Mean.globally());
```


What / Where / When / How

Correctness

Power


Composability

Flexibility


Modularity




1. Classic Batch



2. Batch with Fixed Windows




3. Streaming



4. Streaming with Speculative + Late Data



5. Streaming With Retractions



6. Sessions

What / Where / When / How

Correctness

Power

Composability

Flexibility

Modularity




```
PCollection<KV<String, Integer>> scores = input
    .apply(Sum.integersPerKey());
```

1. Classic Batch

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2))))
    .apply(Sum.integersPerKey());
```

2. Batch with Fixed Windows

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2)))
        .triggering(AtWatermark()))
    .apply(Sum.integersPerKey());
```

3. Streaming

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2)))
        .triggering(AtWatermark())
        .withEarlyFirings(AtPeriod(Minutes(1)))
        .withLateFirings(AtCount(1)))
    .apply(Sum.integersPerKey());
```

4. Streaming with Speculative + Late Data

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Minutes(2)))
        .triggering(AtWatermark())
        .withEarlyFirings(AtPeriod(Minutes(1)))
        .withLateFirings(AtCount(1)))
        .accumulatingAndRetractingFiredPanels())
    .apply(Sum.integersPerKey());
```

5. Streaming With Retractions

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(Sessions.withGapDuration(Minutes(2)))
        .triggering(AtWatermark())
        .withEarlyFirings(AtPeriod(Minutes(1)))
        .withLateFirings(AtCount(1)))
        .accumulatingAndRetractingFiredPanels())
    .apply(Sum.integersPerKey());
```

6. Sessions

What / **Where** / **When** / **How**

Correctness

Power

Composability

Flexibility

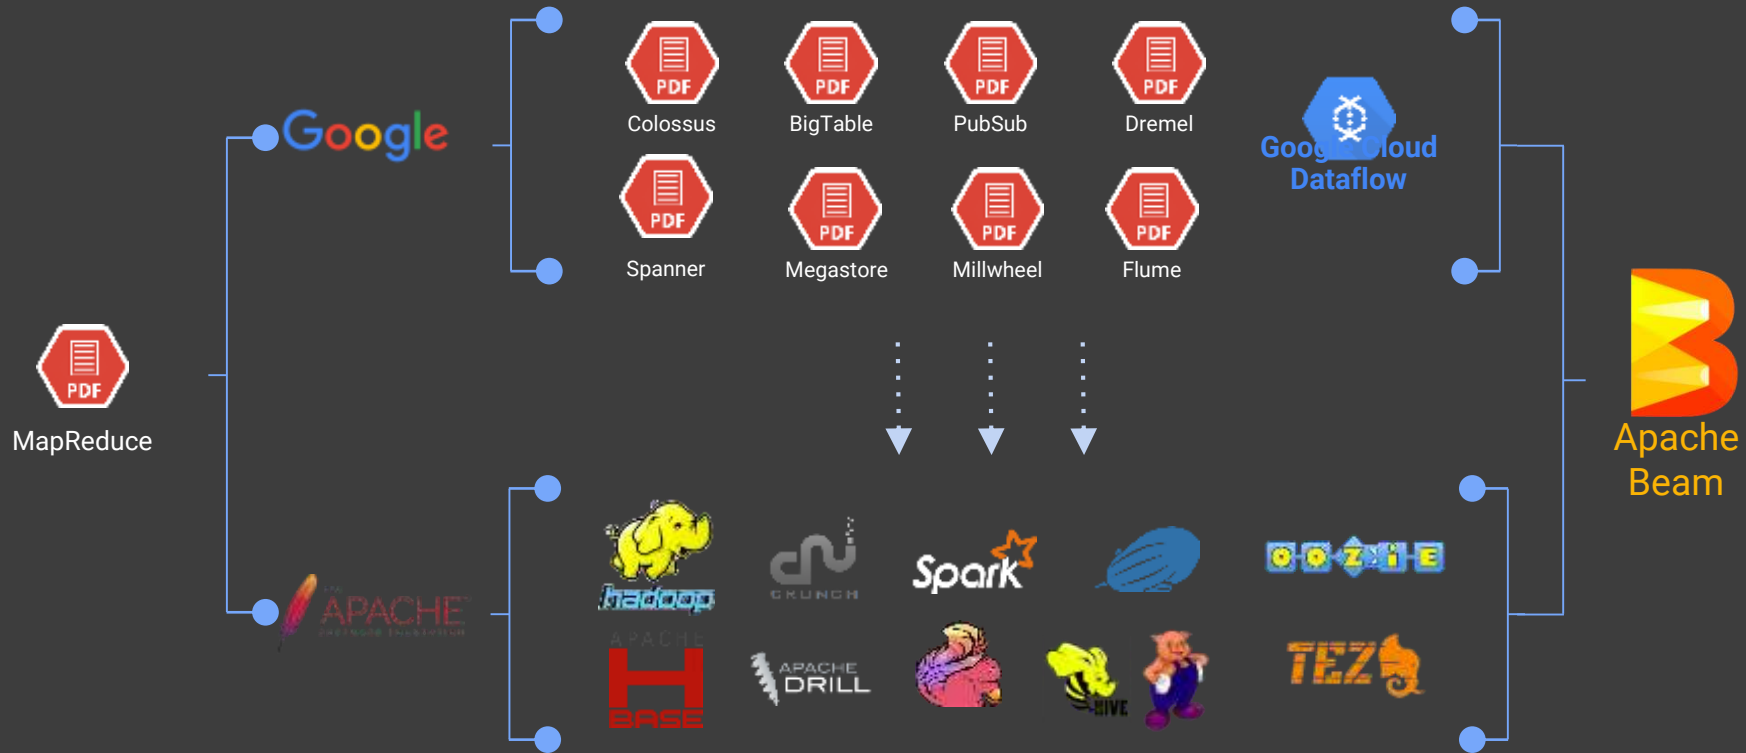
Modularity





Apache Beam (incubating)

The Evolution of Beam



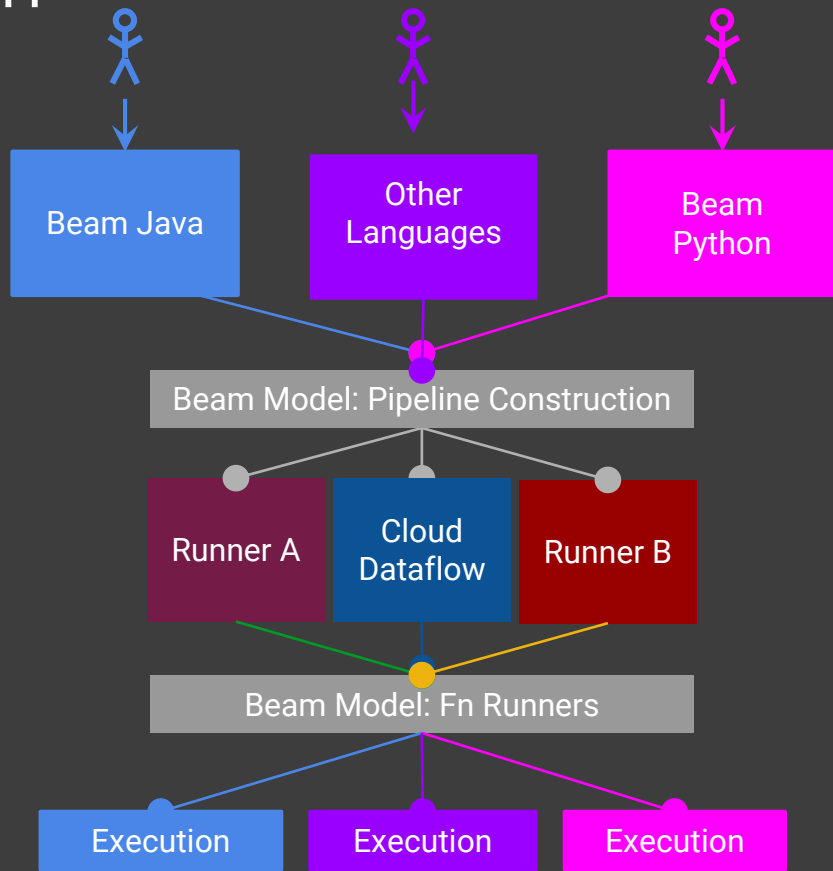
What is Part of Apache Beam?

1. The Beam Model: **What** / **Where** / **When** / **How**
2. SDKs for writing Beam pipelines -- starting with Java
3. Runners for Existing Distributed Processing Backends
 - Apache Flink (thanks to data Artisans)
 - Apache Spark (thanks to Cloudera)
 - Google Cloud Dataflow (fully managed service)
 - Local (in-process) runner for testing



Apache Beam Technical Vision

1. **End users:** who want to write pipelines or transform libraries in a language that's familiar.
2. **SDK writers:** who want to make Beam concepts available in new languages.
3. **Runner writers:** who have a distributed processing environment and want to support Beam pipelines



Categorizing Runner Capabilities

What is being computed?

	Beam Model	Cloud Dataflow	Apache Flink	Apache Spark
ParDo	✓	✓	✓	✓
GroupByKey	✓	✓	✓	~
Flatten	✓	✓	✓	✓
Combine	✓	✓	✓	✓
Composite Transforms	✓	~	~	~
Side Inputs	✓	✓	~	~
Source API	✓	✓	~	✓
Aggregators	~	~	~	~
Keyed State	×	×	×	×

Where in event time?

	Beam Model	Cloud Dataflow	Apache Flink	Apache Spark
Global windows	✓	✓	✓	✓
Fixed windows	✓	✓	✓	~
Sliding windows	✓	✓	✓	×
Session windows	✓	✓	✓	×
Custom windows	✓	✓	✓	×
Custom merging windows	✓	✓	✓	×
Timestamp control	✓	✓	✓	×

When in processing time?

	Beam Model	Cloud Dataflow	Apache Flink	Apache Spark
Configurable triggering	✓	✓	✓	×
Event-time triggers	✓	✓	✓	×
Processing-time triggers	✓	✓	✓	✓
Count triggers	✓	✓	✓	×
[Meta]data driven triggers	×	×	×	×
Composite triggers	✓	✓	✓	×
Allowed lateness	✓	✓	✓	×
Timers	×	×	×	×

How do refinements relate?

	Beam Model	Cloud Dataflow	Apache Flink	Apache Spark
Discarding	✓	✓	✓	✓
Accumulating	✓	✓	✓	×
Accumulating & Retracting	×	×	×	×

<http://beam.incubator.apache.org/capability-matrix/>